

JavaGAT and Ibis Demo



SP 3.1

high-performance distributed computing

***Henri Bal
Niels Drost
Ceriël Jacobs
Jason Maassen
Rob van Nieuwpoort***



www.cs.vu.nl/ibis



vl·e

**1. Develop application
on local workstation**



**2. Submit, monitor and
steer using JavaGAT**




**3a. Deploy on conventional
Grids using Globus**



**3b. Peer-to-Peer deployment
using Zorilla**



4. Application on the Grid!



**1. Develop application
on local workstation**



**2. Submit, monitor and
steer using JavaGAT**




**3a. Deploy on conventional
Grids using Globus**



**3b. Peer-to-Peer deployment
using Zorilla**



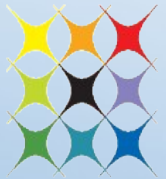
4. Application on the Grid!



**Ibis
programming
models
(Satin)**

Satin

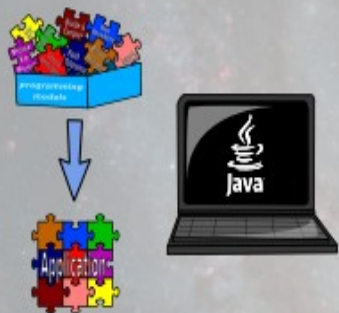
- One of the high-level programming models for the Ibis communication library
- Provide a powerful programming model
 - master/worker, divide-and-conquer, shared objects
- Extremely easy to use
- Satin allows applications to **transparently** deal with grid issues
 - load balancing, malleability, migration, fault-tolerance, adaptivity, firewalls, heterogeneity



vl·e



**1. Develop application
on local workstation**



**2. Submit, monitor and
steer using JavaGAT**




**3a. Deploy on conventional
Grids using Globus**



**3b. Peer-to-Peer deployment
using Zorilla**



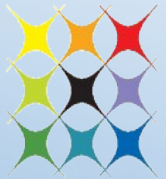
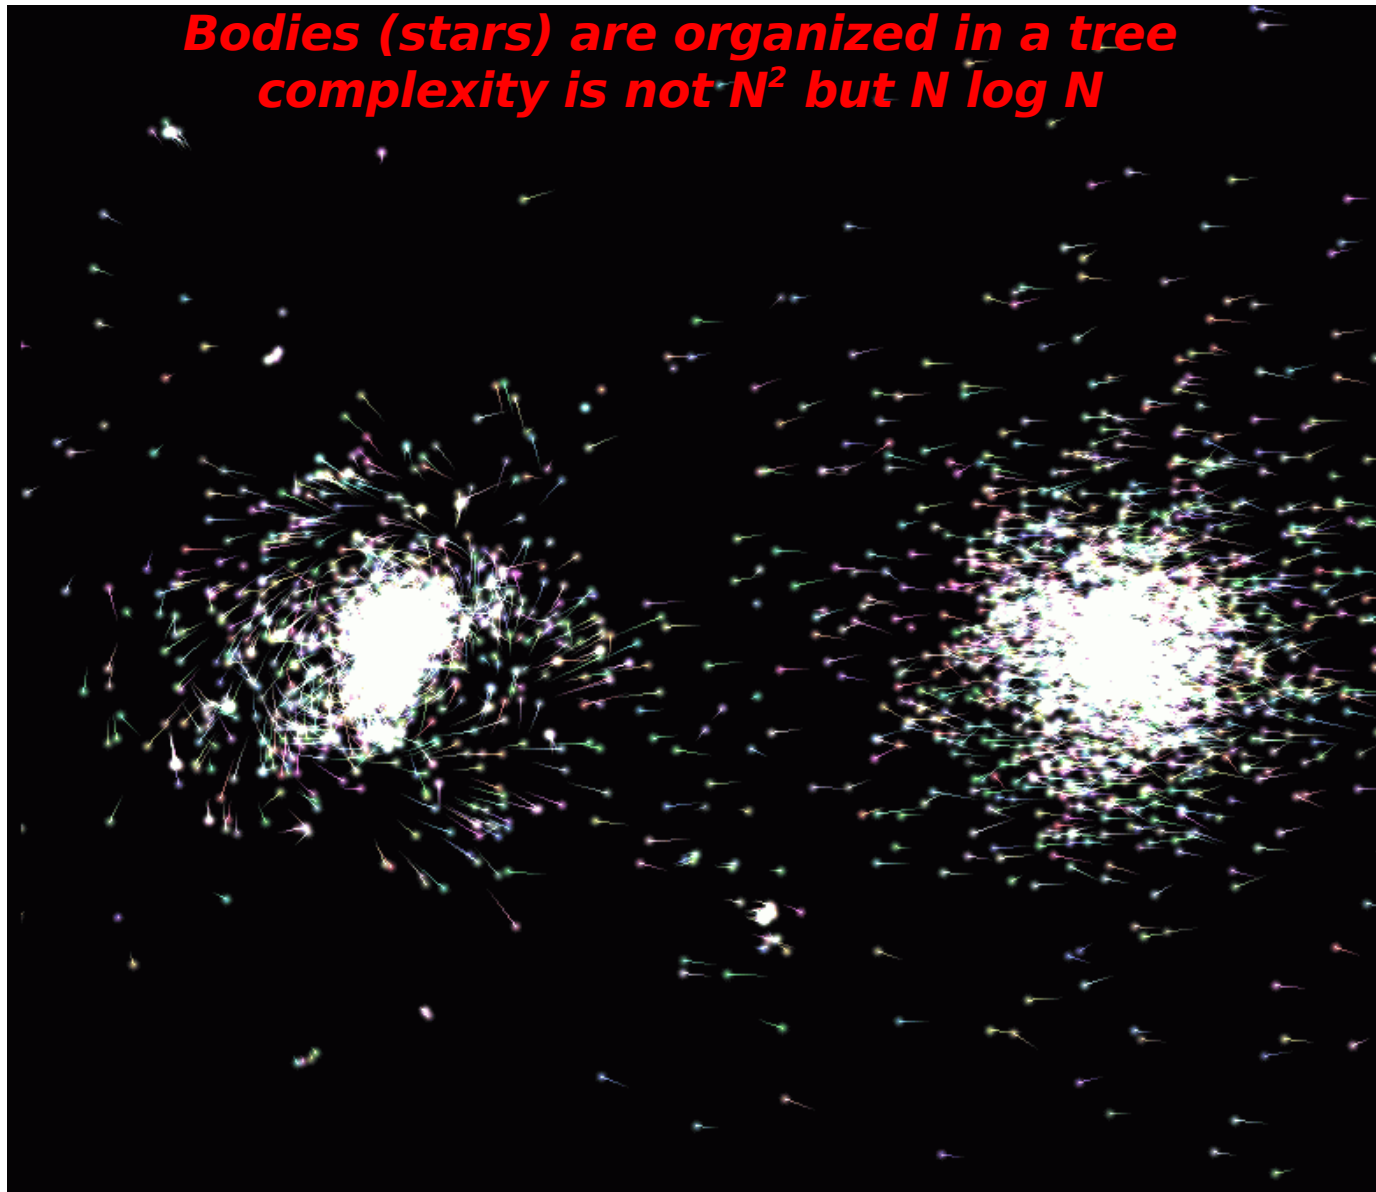
4. Application on the Grid!



**Barnes-Hut
Nbody
simulation**

Barnes-Hut N-body simulation

***Bodies (stars) are organized in a tree
complexity is not N^2 but $N \log N$***



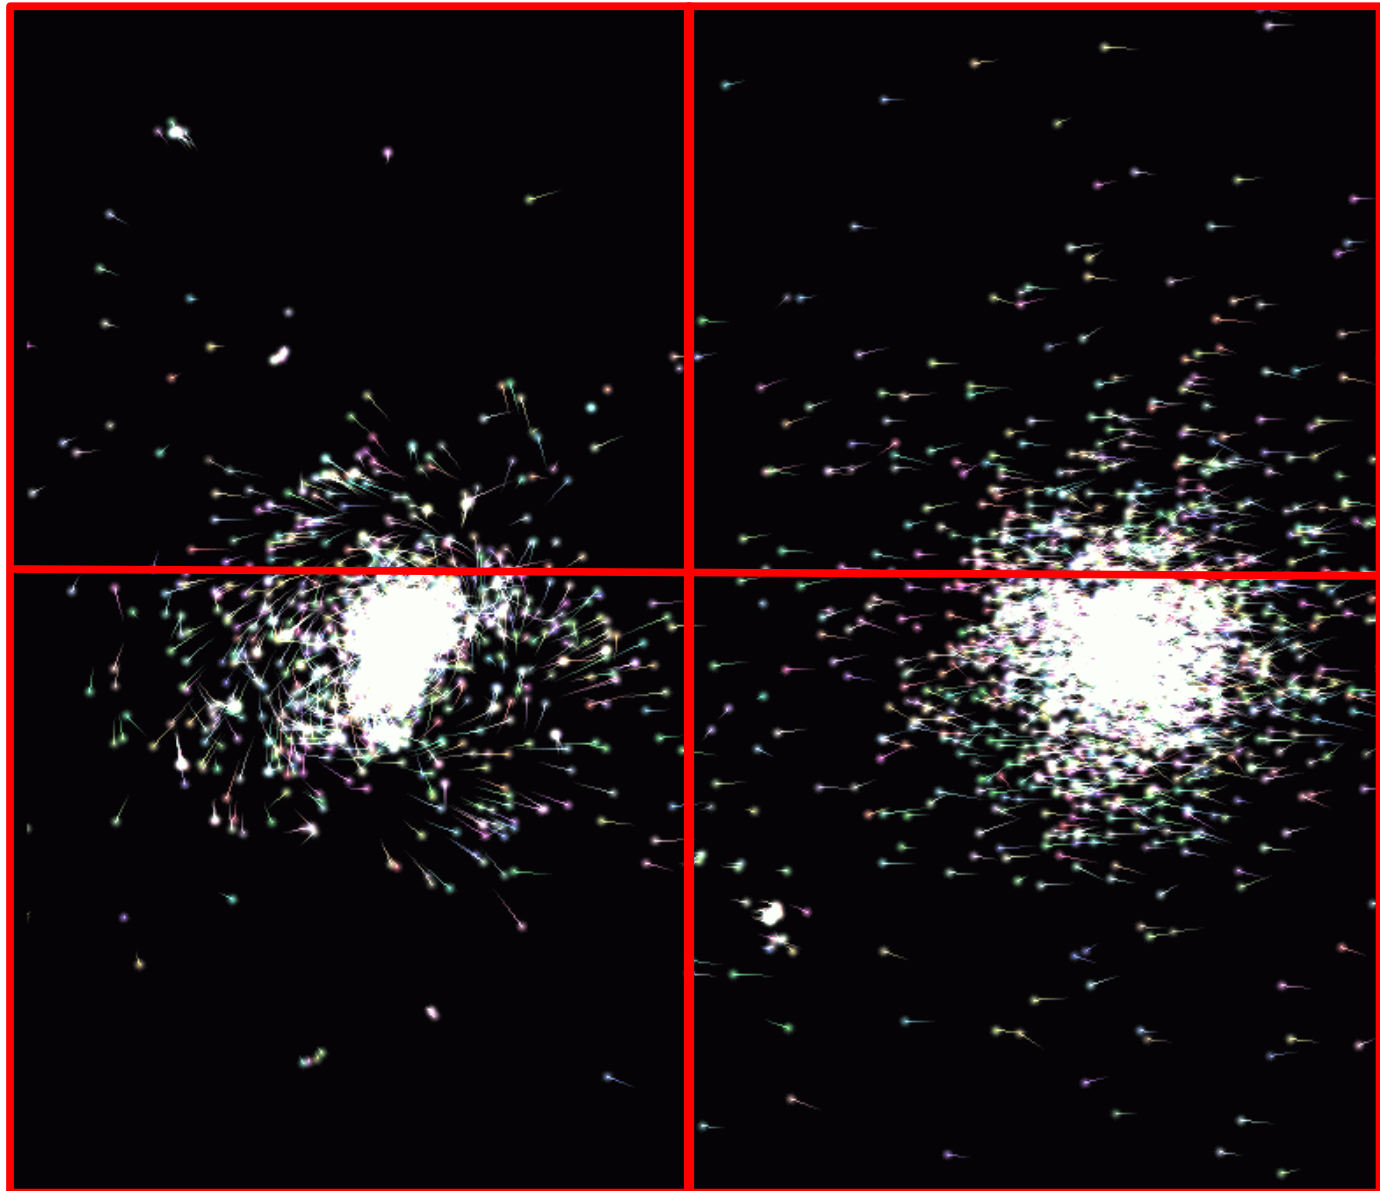
vl·e



ibis



Barnes-Hut N-body simulation



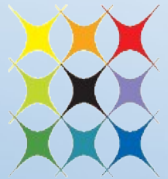
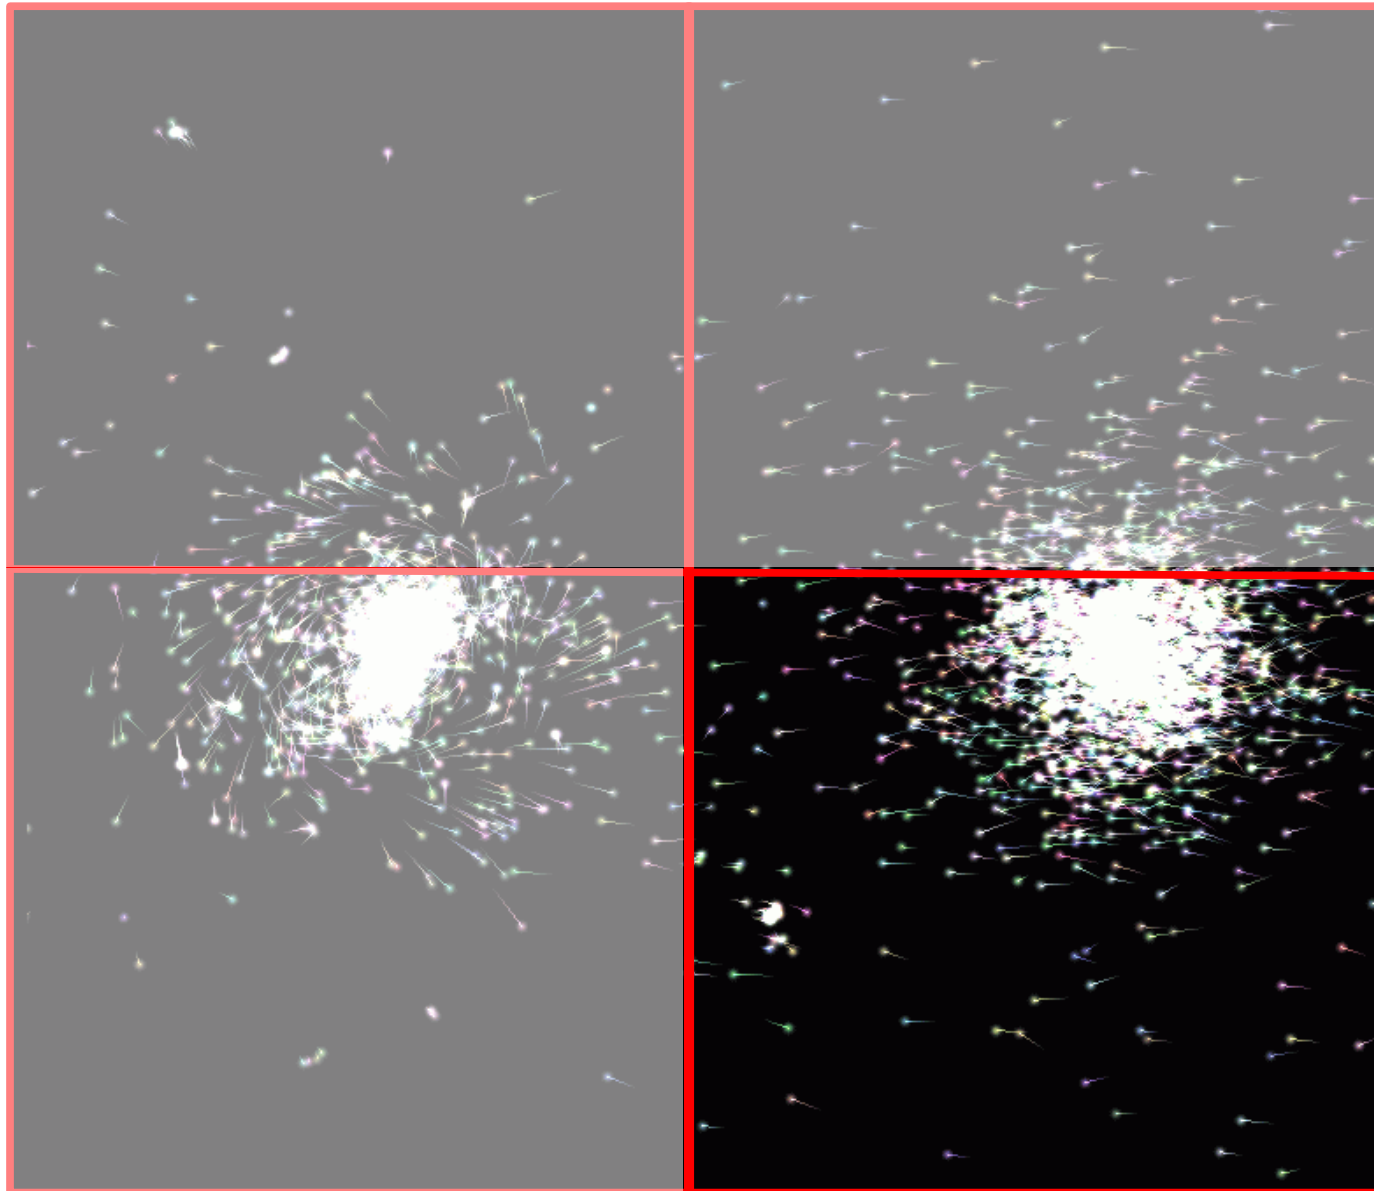
vl·e



ibis



Barnes-Hut N-body simulation





vl·e



ibis



Barnes-Hut N-body simulation



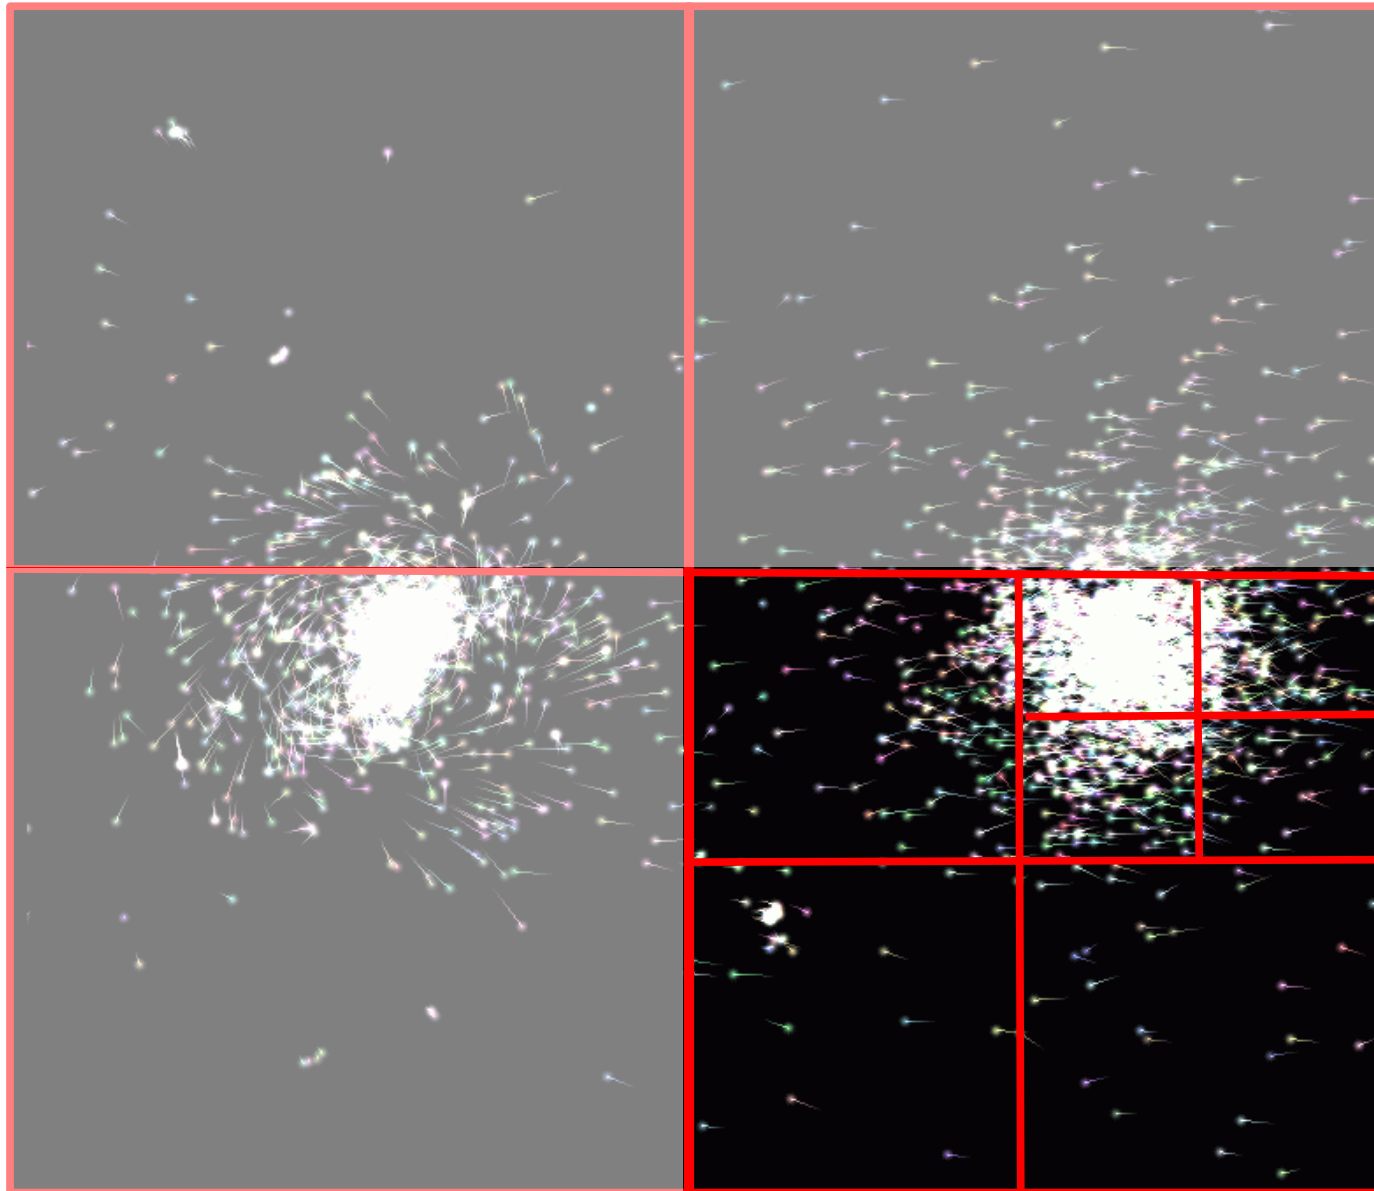
vl·e



ibis



Barnes-Hut N-body simulation



Satin Applications

- VI-e
 - Grammar induction, **SP 2.2**
 - Pieter Adriaans, Cerial Jacobs
 - N-body simulations (DEMO)
- Grammar-based text analysis
- VLSI routing
- Satisfiability solver
- Gene sequencing
- Game-tree search
- Raytracing
- Numerical functions
- ...



**1. Develop application
on local workstation**



**2. Submit, monitor and
steer using JavaGAT**



**Java
Grid
Application
Toolkit
(JavaGAT)**


**3a. Deploy on conventional
Grids using Globus**



**3b. Peer-to-Peer deployment
using Zorilla**




4. Application on the Grid!



The Java Grid Application Toolkit

- API for developing and running **portable** grid applications **independently** of the underlying grid infrastructure and available services
- **Simple** API
- GAT Adaptors (“plugins”) connect GAT to grid services (Globus, Unicore, SSH, ProActive, ...)

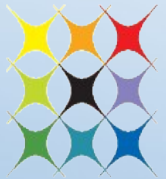


vl-e



Java GAT users

- The Virtual Labs for E-science project (VI-e)
 - AMOLF, Institute for Atomic and Molecular Physics, **SP 1.6**
 - Vrije Universiteit Amsterdam, **SP 3.1**
 - VU Medical Center Amsterdam, **SP 1.3**
 - Vbrowser, **SP 2.5 / SP 1.3**
- The Multimedien project
- Max Planck Institute for Astrophysics in Garching
- D-Grid, Astrogrid
- Louisiana State University
- University of Texas
- The workflow system Triana (University of Cardiff)
- Georgia State University
- Zuse Institute Berlin, Germany
- Download is anonymous, so we don't know



vl·e



vrije Universiteit



Demo ...

Distributed supercomputing

- Parallel processing on geographically distributed computing systems (grids)
- Our goals:
 - Don't use individual supercomputers / clusters, but combine multiple systems
 - Provide high-level programming support



Optimizing for the grid

- Grids usually are **hierarchical**
 - Collections of clusters, supercomputers
 - Fast local links, slow wide-area links
- Optimize algorithms to exploit hierarchy
 - Message combining + latency hiding on wide-area links
 - Collective operations for wide-area systems
 - Successful for many applications



Satin

master-worker

- Master-worker parallelism
 - Divide work into parts
 - Spawn job for each part
 - Solve parts in parallel
 - Combine results



vl·e



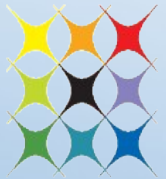
ibis



Satin

divide-and-conquer

- Parallel Divide-and-conquer
 - Divide work into parts
 - Spawn job for each part
 - Solve parts in parallel
 - Combine results
 - **But now recursively!**
 - **sub-problems split the work up further and spawn their own sub-jobs**



vl·e



Satin

- Divide-and-conquer

job1



vl·e




ibis



Satin

- Divide-and-conquer



vl·e

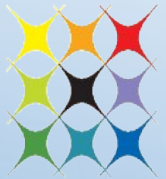



ibis



Satin

- Divide-and-conquer



vl·e





ibis



Satin

- Divide-and-conquer



vl·e




ibis



Satin

- Divide-and-conquer



vl·e




ibis



Satin

- Fits hierarchical structure of Grids



vl·e

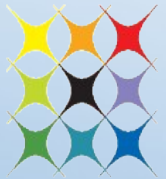


ibis



The Grid

- Satin explicitly targets the grid
- Different architectures
- Firewalls
- Slow networks (latency)
- Distributed memory
- Machines come and go
- Machines crash
- Machines have different speeds



vl·e



ibis



Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;

    void updateBodies(BodyUpdates b, int iter) { // Global method.
        root.applyUpdates(b, iter); // Update bodies in our tree.
    }

    BodyTreeNode getRoot() { // Local method.
        return root;
    }
}


// Mark the computeForces method as a spawn operation.
interface BHSpawns extends satin.Spawnable {
    BodyUpdates computeForces(Subtree s, int iter, Bodies bodies);
}

class BarnesHut extends satin.SatinObject implements BHSpawns {
    boolean guard_computeForces(Subtree s, int iter, Bodies bodies) {
        return bodies.iter + 1 == iter;
    }

    // Spawnable method. The "bodies" parameter is a shared object.
    BodyUpdates computeForces(Subtree s, int iter, Bodies bodies) {
        if(s.hasNoChildren) {
            computeSequentially(s, iter, bodies.getRoot());
        } else { // Divide the work and spawn tasks (recursion step).
            for(int i=0; i<s.nrChildren; i++) {
                res[i] = computeForces(s.child[i], iter, bodies); // Spawn.
            }
            sync(); // Wait for the spawn operation to finish.

            return mergeSubresults(res); // Merge results and return.
        }
    }

    public static void main(String[] args) {
        BarnesHut bh = new BarnesHut();
        Bodies bodies = new Bodies(); // Create shared object.
        for (int iter = 0; iter < N; iter++) {
            results = bh.computeForces(root, iter, bodies); // Spawn.
            sync(); // Wait for the spawn operation to finish.
            bodies.update(results, iter); // Shared method invocation.
        }
    }
}
```



vl·e



ibis



Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;

    void updateBodies(BodyUpdates b, int iter) { // Global method.
        root.applyUpdates(b, iter); // Update bodies in our tree.
    }

    BodyTreeNode getRoot() { // Local method

```

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}
```

```
// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;
```

```
void updateBodies(BodyUpdates b, int iter) {
    root.applyUpdates(b, iter);
}
```

```
BodyTreeNode getRoot() {
    return root;
}
```

```
}
```


Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;

    void updateBodies(BodyUpdates b, int iter) { // Global method.
        root.applyUpdates(b, iter); // Update bodies in our tree.
    }

    BodyTreeNode getRoot() { // Local method.
        return root;
    }
}

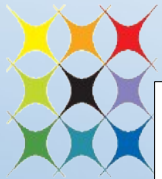
// Mark the computeForces method as a spawn operation.
interface BHSpawns extends satin.Spawnable {
    BodyUpdates computeForces(Subtree s, int iter, Bodies bodies);
}

class BarnesHut extends satin.SatinObject implements BHSpawns {
    boolean guard computeForces(Subtree s, int iter, Bodies bodies) {
```

```
// Mark the computeForces method as a spawn operation.
interface BHSpawns extends satin.Spawnable {
    BodyUpdates computeForces(Subtree s, int iter, Bodies bodies);
}
```

```
    return mergeSubresults(res); // Merge results and return.
}

public static void main(String[] args) {
    BarnesHut bh = new BarnesHut();
    Bodies bodies = new Bodies(); // Create shared object.
    for (int iter = 0; iter < N; iter++) {
        results = bh.computeForces(root, iter, bodies); // Spawn.
        sync(); // Wait for the spawn operation to finish.
        bodies.update(results, iter); // Shared method invocation.
    }
}
```



vl·e



ibis



Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;
```

```
public static void main(String[] args) {
    BarnesHut bh = new BarnesHut();
    Bodies bodies = new Bodies();
    for (int iter = 0; iter < N; iter++) {
        results = bh.computeForces(root, iter, bodies);
        sync();
        bodies.updateBodies(results, iter);
    }
}
```

```
    results = bh.computeForces(root, iter, bodies); // Spawn.
    sync(); // Wait for the spawn operation to finish.
    return mergeSubresults(res); // Merge results and return.
}

public static void main(String[] args) {
    BarnesHut bh = new BarnesHut();
    Bodies bodies = new Bodies(); // Create shared object.
    for (int iter = 0; iter < N; iter++) {
        results = bh.computeForces(root, iter, bodies); // Spawn.
        sync(); // Wait for the spawn operation to finish.
        bodies.updateBodies(results, iter); // Shared method invocation.
    }
}
```



vl.

}



ibis



Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;

    void updateBodies(BodyUpdates b, int iter) { // Global method
```

```
BodyUpdates computeForces(Subtree s, int iter, Bodies bodies) {
    if(s.hasNoChildren) {
        computeSequentially(s, iter, bodies.getRoot());
    } else {
        for(int i=0; i<s.nrChildren; i++) {
            res[i] = computeForces(s.child[i], iter, bodies);
        }
        sync();

        return mergeSubresults(res);
    }
}
```

```
public static void main(String[] args) {
    BarnesHut bh = new BarnesHut();
    Bodies bodies = new Bodies(); // Create shared object.
    for (int iter = 0; iter < N; iter++) {
        results = bh.computeForces(root, iter, bodies); // Spawn.
        sync(); // Wait for the spawn operation to finish.
        bodies.update(results, iter); // Shared method invocation.
    }
}
```



ibis



Barnes-Hut Nbody code

```
// Marker interface that defines updateBodies as a global method.
interface BodiesInterface extends satin.GlobalMethods {
    void updateBodies(BodyUpdates b, int iter);
}

// A shared object containing the tree of bodies.
class Bodies extends satin.SharedObject implements BodiesInterface {
    BodyTreeNode root;

    void updateBodies(BodyUpdates b, int iter) { // Global method.
        root.updateBodies(b, iter); // Update bodies in our tree
    }
}
```

```
boolean guard_computeForces(Subtree s, int iter, Bodies bodies) {
    return iter == bodies.iter + 1;
}
```

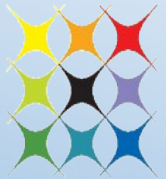
```
}

class BarnesHut extends satin.SatinObject implements BHSpawns {
    boolean guard_computeForces(Subtree s, int iter, Bodies bodies) {
        return bodies.iter + 1 == iter;
    }

    // Spawnable method. The "bodies" parameter is a shared object.
    BodyUpdates computeForces(Subtree s, int iter, Bodies bodies) {
        if(s.hasNoChildren) {
            computeSequentially(s, iter, bodies.getRoot());
        } else { // Divide the work and spawn tasks (recursion step).
            for(int i=0; i<s.nrChildren; i++) {
                res[i] = computeForces(s.child[i], iter, bodies); // Spawn.
            }
            sync(); // Wait for the spawn operation to finish.

            return mergeSubresults(res); // Merge results and return.
        }
    }

    public static void main(String[] args) {
        BarnesHut bh = new BarnesHut();
        Bodies bodies = new Bodies(); // Create shared object.
        for (int iter = 0; iter < N; iter++) {
            results = bh.computeForces(root, iter, bodies); // Spawn.
            sync(); // Wait for the spawn operation to finish.
            bodies.update(results, iter); // Shared method invocation.
        }
    }
}
```



vl·e

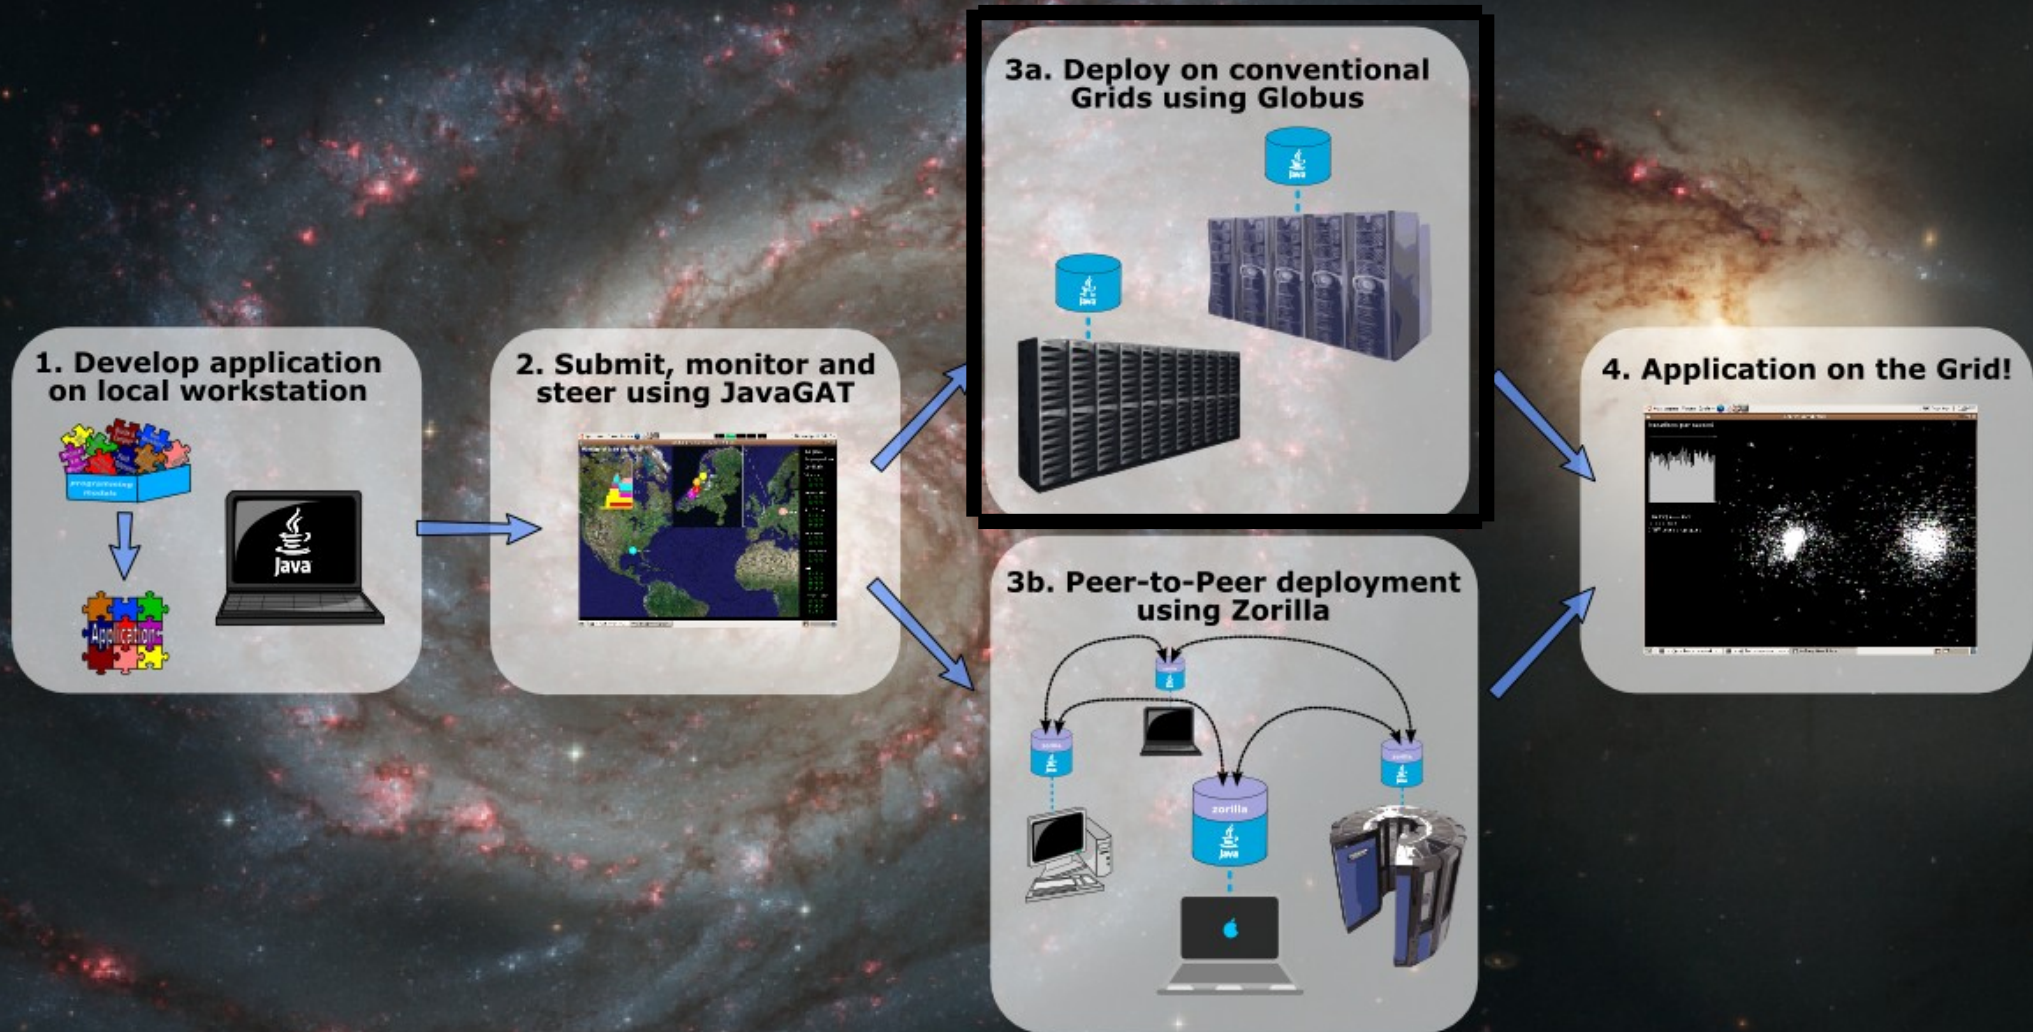


ibis

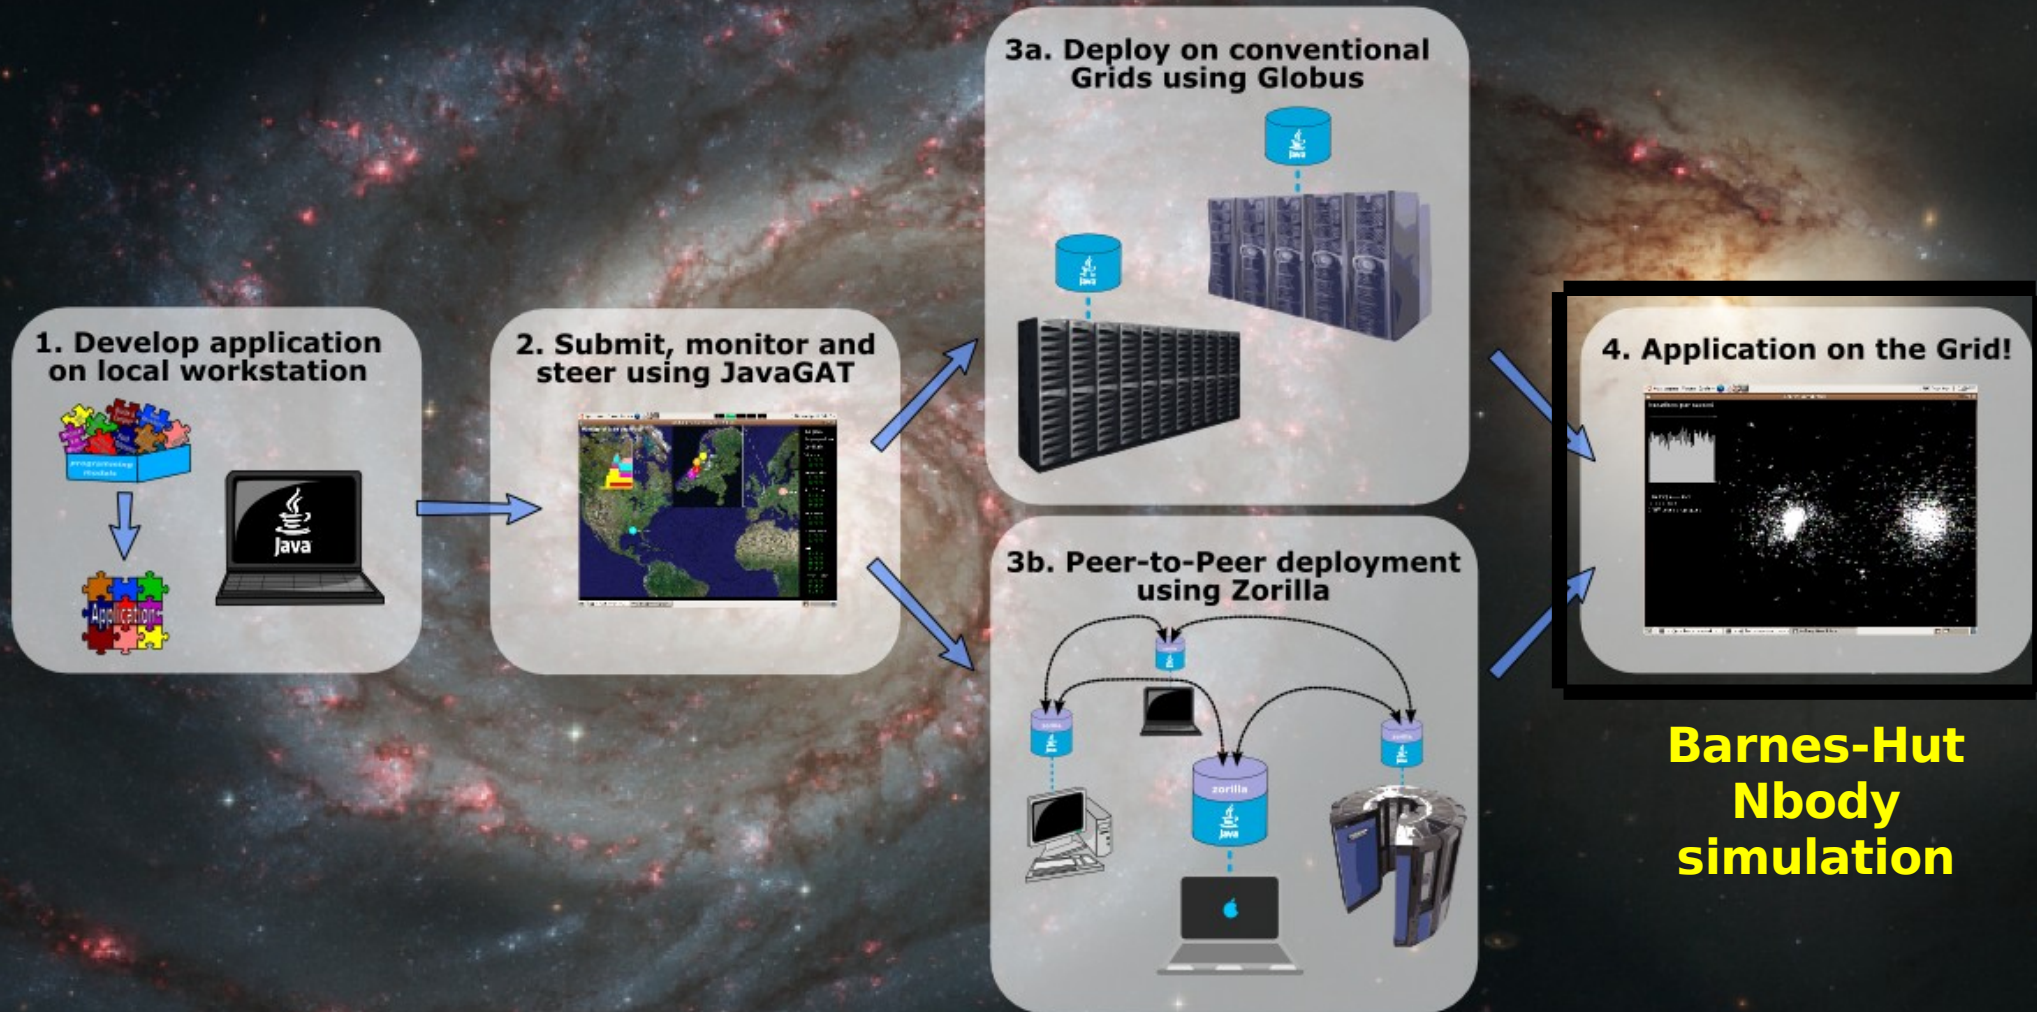


Application Example

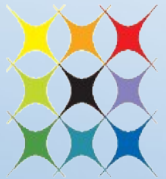
Grid middleware (Globus)



Application Example



Sequential Fibonacci



vl·e



ibis

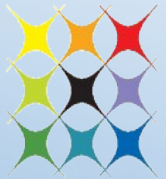


```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
  
    return x + y;  
}
```

Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
    sync();  
    return x + y;  
}
```



vl·e



ibis

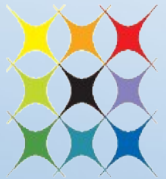


Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;  
  
    long x = fib(n - 1);  
    long y = fib(n - 2);  
    sync();  
    return x + y;  
}
```

Mark methods as
Spawnable.
They can run in parallel.



vl·e



ibis



Parallel Fibonacci

```
interface FibInterface extends ibis.satin.Spawnable {  
    public long fib(int n);  
}
```

```
public long fib(int n) {  
    if (n < 2) return n;
```

```
    long x = fib(n - 1);
```

```
    long y = fib(n - 2);
```

```
    sync();
```

```
    return x + y;
```

```
}
```

Mark methods as

Spawnable.

They can run in parallel.

Wait until spawned
methods are done.



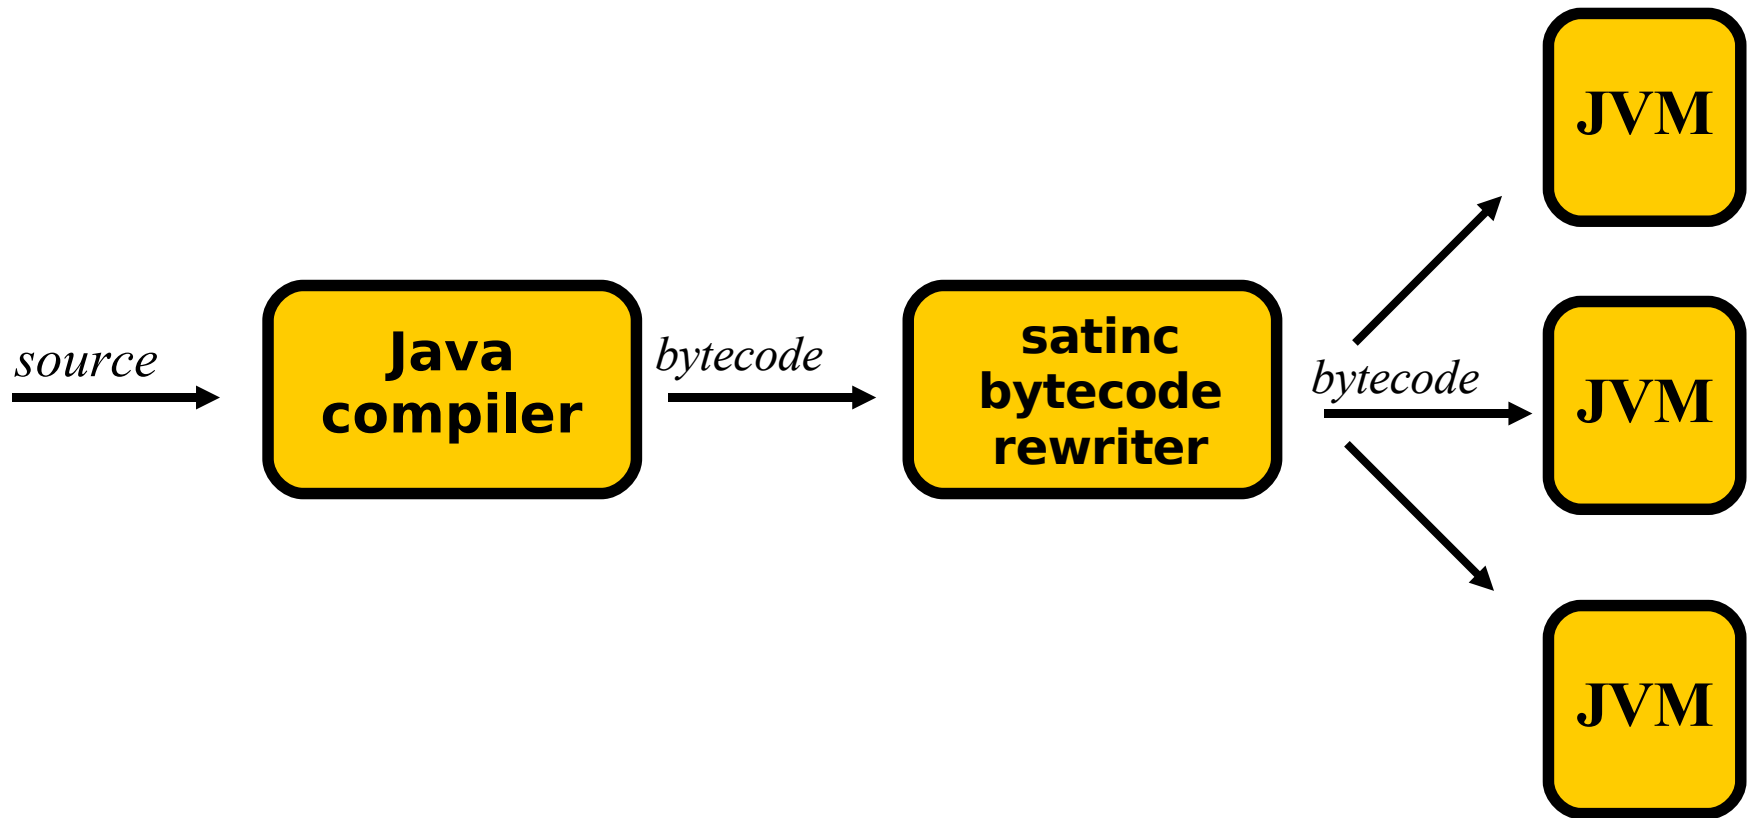
vl·e



ibis



Compiling Satin Programs



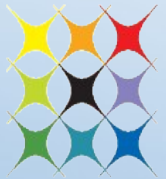
Java GAT API features

- Security (deal with passwords, credentials, etc)
- Grid I/O
 - File operations, remote file access, file replication
 - Inter-process communication
- Resource Management
 - Resource brokering
 - Forking grid applications, job management
- Application Information Management
 - Global repository for application specific information
 - Query this information repository
- Monitoring
 - Grid monitoring
 - Application monitoring and steering



Java GAT File example

```
public class RemoteCopy {  
    public static void main(String[] args) {  
        GATContext context = new GATContext();  
  
        URI src = new URI(args[0]);  
        URI dest = new URI(args[1]);  
        File file = GAT.createFile(context, src);  
        file.copy(dest);  
        GAT.end();  
    }  
}
```



vl·e



ibis

